

# **Regular Expressions**

Genome 559: Introduction to Statistical and  
Computational Genomics

**Elhanan Borenstein**


# A quick review: The super *Date* class

```
class Date:
    def __init__(self, day, month):
        self.day = day
        self.month = month
    def __str__(self) :
        day_str = '%s' % self.day
        mon_str = self.month
        return mon_str + "-" + day_str

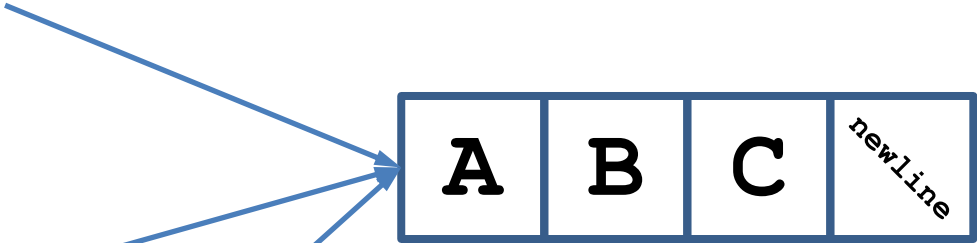
birthday = Date(3,"Sep")
print "It's ", birthday, ". Happy Birthday!"
```


```
It's Sep-3. Happy Birthday!
```


# Strings

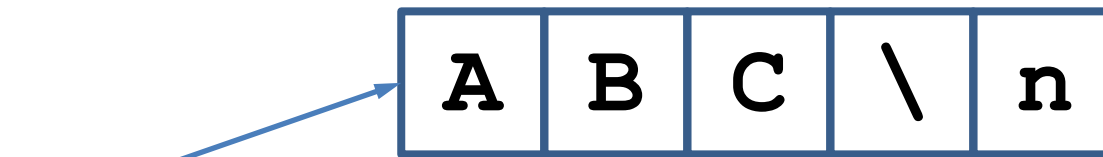


# Newlines are a bit more complicated

- `'abc\n'`

<b>A</b>	<b>B</b>	<b>C</b>	<i>newline</i>
----------	----------	----------	----------------
- `"abc\n"`

<b>A</b>	<b>B</b>	<b>C</b>	<i>newline</i>
----------	----------	----------	----------------
- `"""abc  
"""`

<b>A</b>	<b>B</b>	<b>C</b>	<i>newline</i>
----------	----------	----------	----------------
- `r'abc\n'`

<b>A</b>	<b>B</b>	<b>C</b>	<b>\</b>	<b>n</b>
----------	----------	----------	----------	----------

# Why so many?

- ‘ vs “ lets you put the other kind inside a string. Very Useful.
- ''' lets you run across multiple lines.
- All 3 let you include and show *invisible* characters (using \n, \t, etc.)
- r'...' (*raw strings*) do not support invisible character, but avoid problems with backslash. Will become useful very soon.

```
open('C:\new\text.dat') vs.  
open('C:\\new\\text.dat') vs.  
open(r'C:\new\text.dat')
```

# String operations

- As you recall, the string data type supports a variety of operations:

```
>>> my_str = 'tea for too'
>>> print my_str.replace('too', 'two')
'tea for two'

>>> print my_str.upper()
TEA FOR TOO

>>> my_str.split(' ')
['tea', 'for', 'too']

>>> print my_str.find("o")
5

>>> print my_str.count("o")
3
```

# But ...

- What if we want to do more complex things?
  - Get rid of all punctuation marks
  - Find all dates in a long text and convert them to a specific format
  - Delete duplicated words
  - Find **all** email addresses in a long text
  - Find everything that “looks” like a gene name in some output file
  - Split a string whenever a certain word (rather than a certain character) occurs
  - Find DNA motifs in a Fasta file



# Well ...

- We can always write a program that does that ...

```
# assume we have a genome sequence in string variable myDNA
for index in range(0,len(myDNA)-20) :
    if (myDNA[index] == "A" or myDNA[index] == "G") and
        (myDNA[index+1] == "A" or myDNA[index+1] == "G") and
        (myDNA[index+2] == "A" or myDNA[index+2] == "G") and
        (myDNA[index+3] == "C") and
        (myDNA[index+4] == "A") and
        # and on and on!
    ...
    (myDNA[index+19] == "C" or myDNA[index+19] == "T") :
        print "Match found at ",index
        break
```





# Regular expressions

- Regular expressions (a.k.a. RE, regexp, regexes, regex) are a highly specialized **text-matching tool**.
- They are extremely useful in searching and modifying (long) string
- Regex can be viewed as a tiny programming language embedded in Python and made available through the re module.
- <http://docs.python.org/library/re.html>

# Not only in Python

- REs are very widespread:
  - Unix utility “grep”
  - Perl
  - TextWrangler
  - TextPad
  - Python
- So, ... learning the “RE language” would serve you in many different environments as well.

# Do you absolutely need regexes?

- No, everything they do, you could do yourself!
- BUT ... pattern-matching is:
  - Widely used (especially in bioinf applications)!
  - Tedious to program!
  - Error-prone!
- RE give you a flexible, systematic, compact, and automatic way to do it.  
(In truth, it's still somewhat error-prone, but in a different way).

# RE is It's all about finding a great match

- Using this RE tiny language, you can specify patterns that you want to match
- You can then ask ***match*** questions such as:
  - “Does this string **match** this **pattern**?”
  - “Is there a **match** to this **pattern** anywhere in this string?”
  - “What are all the **matches** to this **pattern** in this string?”
- You can also use REs to **modify** a string
  - **Replace** parts of a string (sub) that **match** the **pattern** with something else
  - Break strings into smaller pieces (split) wherever this **pattern** is **matched**

# A simple example

- Consider the following example:

```
>>> import re
>>> re.findall(r'\bf[a-z]*', 'which foot or hand fell fastest')
['foot', 'fell', 'fastest']
```

This RE means: A word that starts with 'f' followed by any number of alphabetical characters

- Note the *re.* prefix – findall is a function in the re module
- findall:
  - Format: `findall(<regexe>, <string>)`
  - Returns a list of all non-overlapping substrings that matches the regexe.
- REs are provided as strings.

Remember:  
It's all about matching

*Regular expressions are patterns;  
they “match” sequences of characters*

# Basic RE matching

- Most letters and numbers match themselves
  - For example, the regular expression `test` will match the string `test` exactly
  - Normally case sensitive

```
>>> re.findall(r'test', "Tests are testers' best testimonials")  
['test', 'test']
```

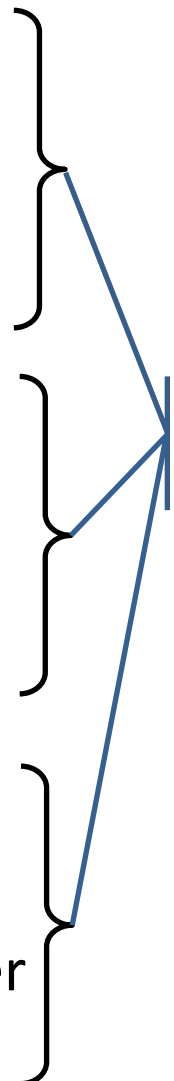
- Most punctuation marks have special meanings!
  - Metacharacters: `.` `^` `$` `*` `+` `?` `{` `[` `]` `\` `|` `(` `)`
  - needs to be escaped by backslash (e.g., `\"` instead of `.`) to get non-special behavior
  - Therefore, “raw” string literals (`r'C:\new.txt'`) are generally recommended for regexes (unless you double your backslashes judiciously)

# Sets

- **Square brackets** mean that any of the listed characters will do (matching one of several alternatives)
  - `[abc]` means either "a" , "b" , or "c"
- You can also give a range:
  - `[a-d]` means "a", "b", "c", or "d"
- Negation: caret means ***not***
  - `[^a-d]` means anything but a, b, c or d
  - `[^5]` means anything but 5
- Metacharacters are not active inside sets.
  - `[ak$]` will match "a", "k", or "\$". Normally, "\$" is a metacharacter. Inside a set it's stripped of its special nature.



# Predefined sets

- `\d` matches any decimal digit (equivalent to `[0-9]`).
  - `\D` matches any non-digit character (equivalent to `[^0-9]`).
  - `\s` matches any whitespace character (equivalent to `[\t\n\r\f\v]`).
  - `\S` matches any non-whitespace character (equivalent to `[^\t\n\r\f\v]`).
  - `\w` matches any alphanumeric character (equivalent to `[a-zA-Z0-9_]`).
  - `\W` matches any non-alphanumeric character (equivalent to the class `[^a-zA-Z0-9_]`).
- 

Note the pairs.  
Easy to remember!

# Matching boundaries

- `^` matches the beginning of the string
- `$` matches the end of the string
- `\b` matches a word boundary
- `\B` matches position that is not a word boundary

(A word boundary is a position that changes from a word character to a non-word character, or vice versa).

For example, `\bcat` will match **catalyst** but not **location**

# Wildcards

- `.` matches **any** character (except newline)
- If you really mean “.” you must use a backslash
- WARNING:
  - backslash is special in Python strings
  - It's special again in RE
  - This means you need too many backslashes
  - Use “raw strings” to make things simpler
- What does this RE means: `r'\d\.\d'?`

# Repetitions

- Allows you to specify that a portion of the RE must/can be repeated a certain number of times.
- **\*** : The previous character can repeat 0 or more times
  - `ca*t` matches "ct", "cat", "caat", "caaat" etc.
- **+** : The previous character can repeat 1 or more times
  - `ca+t` matches "cat", "caat" etc. but not "ct"
- Braces provide a more detailed way to indicate repeats
  - `A{1, 3}` means at least one and no more than three A's
  - `A{4, 4}` means exactly four A's

# A quick example

- Remember this PSSM:



```
re.findall(r'[AG]{3,3}CATG[TC]{4,4}[AG]{2,2}C[AT]TG[CT][CG][TC]', myDNA)
```

# More examples

```
>>> re.sub('\d', 'x', 'a_b - 12')
'a_b - xx'
>>> re.sub('\D', 'x', 'a_b - 12')
'xxxxxx12'
>>> re.sub('\s', 'x', 'a_b - 12')
'a_bx-x12'
>>> re.sub('\S', 'x', 'a_b - 12')
'xxx x xx'
>>> re.sub('\w', 'x', 'a_b - 12')
'xxx - xx'
>>> re.sub('\W', 'x', 'a_b - 12')
'a_bxxx12'
>>> re.sub('^', 'x', 'a_b - 12')
'xa_b - 12'
>>> re.sub('$', 'x', 'a_b - 12')
'a_b - 12x'
>>> re.sub('\b', 'x', 'a_b - 12')
'a_b - 12'
>>> re.sub('\b\b', 'x', 'a_b - 12')
'xa_bx - x12x'
>>> re.sub(r'\b', 'x', 'a_b - 12')
'xa_bx - x12x'
>>> re.sub('\B', 'x', 'a_b - 12')
'ax_xb x-x 1x2'
```

# RE Semantics

- If  $R, S$  are regexes:
  - $RS$  matches the concatenation of strings matched by  $R, S$  individually
  - $R|S$  matches the union (either  $R$  or  $S$ )
  - `this|that` matches 'this' and 'that', but not 'thisthat'.
- Parentheses can be used for grouping
  - `(abc) +` matches 'abc', 'abcabc', 'abcabcabc', etc.

# Conflicts?

- Check this example:

```
>>> import re
>>> mystring = "This contains 2 files, hw3.py and uppercase.py."
>>> all_matches = re.findall(r' .+\.py', mystring)
>>> print all_matches
```

- What do you think all\_matches contains?

```
[' This contains 2 files, hw3.py and uppercase.py']
```

## What happened?



# Matching is greedy

```
>>> import re
>>> mystring = "This contains 2 files, hw3.py and uppercase.py."
>>> all_matches = re.findall(r' .+\.py', mystring)
>>> print all_matches
[' This contains 2 files, hw3.py and uppercase.py']
```

- Our RE matches “*hw3.py*”
- Unfortunately ...
  - It also matches: “This contains 2 files, *hw3.py*”
  - And it even matches: “This contains 2 files, *hw3.py* and uppercase.py”
- **Python will choose the longest match!**
- Solution:
  - Break my text first into words (not an ideal solution)
  - I could specify that no spaces are allowed in my match

# A better version

- This will work:

```
>>> import re
>>> mystring = "This contains 2 files, hw3.py and uppercase.py."
>>> all_matches = re.findall(r' [^ ]+\.py', mystring)
>>> print all_matches
```

```
['hw3.py' , 'uppercase.py' ]
```

`r".+\.py"`     "Two files: hw3.py and upper.py."

`r"\w+\.py"`     "Two files: hw3.py and UPPER.py."

# What (else) can we do with RE

- `re.findall(pat, str)`
  - finds all (nonoverlapping) matches
- `re.match(pat, str)`
  - matches only at the beginning of the string
- `re.search(pat, str)`
  - matches anywhere in the string
- And also split and substitute...

# What do these functions return

- `re.findall(pat, str)`
  - finds all (nonoverlapping) matches

If nothing was found:  
returns an empty list  
Otherwise:  
returns a list of  
strings

- `re.match(pat, str)`
  - matches only at the beginning of the string

If nothing was found:  
returns None  
Otherwise:  
returns a

- `re.search(pat, str)`
  - matches anywhere in the string

**“match” object**

- And also split and substitute...

For more on RegExp

*(and to learn how to use RegExp to edit strings)*

see next presentation !!

# RE Quick Reference

## MATCHING CHARACTER SETS

- Most letters and numbers match themselves
- `[abc]` means either "a", "b", or "c"
- `[a-d]` means "a", "b", "c", or "d"
- `[^a-d]` means anything but a, b, c or d
- `\d` matches any decimal digit (equivalent to `[0-9]`).
- `\D` matches any non-digit character (equivalent to `[^0-9]`).
- `\s` matches any whitespace character (equivalent to `[\t\n\r\f\v]`).
- `\S` matches any non-whitespace character (equivalent to `[^\t\n\r\f\v]`).
- `\w` matches any alphanumeric character (equivalent to `[a-zA-Z0-9_]`).
- `\W` matches any non-alphanumeric character (equivalent to the class `[^a-zA-Z0-9_]`).
- `.` matches **any** character (except newline)

## MATCHING BOUNDARIES

- `^` matches the beginning of the string
- `$` matches the end of the string
- `\b` matches a word boundary
- `\B` matches position that is not a word boundary

## REPETITION

- `*` : The previous character can repeat 0 or more times
- `+` : The previous character can repeat 1 or more times
- `A{1, 3}` means at least one and no more than three A's

## SEMANTICS

- `RS` matches the concatenation of strings matched by R, S individually
- `R|S` matches the union (either R or S)

## RE FUNCTIONS/PATTERN OBJECT METHODS

- `re.findall(pat, str)`  
Finds all (non-overlapping) matches
- `re.match(pat, str)`  
Matches only at the beginning of str
- `re.search(pat, str)`  
Matches anywhere in str
- `re.split(pat, str)`  
Splits str anywhere matches are found
- `re.sub(pat, new_str, str)`  
Substitutes matched patterns in str with new\_str
- `re.compile(pat)`  
Compile a Pattern object

## MATCH OBJECT METHODS

- `group()` :  
Returns the string that was matched
- `group(i)` :  
Returns the *i* sub-pattern that was matched
- `groups()` :  
Returns all sub-patterns that were matched as a list
- `start()` :  
Returns starting position of the match
- `end()` :  
Returns ending position of the match
- `span()` :  
Returns (start,end) as a tuple

# Regexe vs. Python

- The regular expression language is relatively small and restricted
  - Not all possible string processing tasks can be done using regular expressions.
  - Some tasks can be done with RE, but the expressions turn out to be **extremely** complicated.
- In these cases, you may be better off writing a Python code to do the processing:
  - Python code may take longer to write
  - It will be slower than an elaborate regular expression
  - But ... it will also probably be more understandable.

# Code like a pro ...



- Suppose you are not sure:
  - ... whether the format you are using for a certain command is the correct one
  - or ... whether `range(4)` returns 0 to 4 or 0 to 3
  - or ... whether string has a method “reverse”
  - or ... whether you are allowed to break inside a nested loop

**What should you do?**





# Code like a pro ...



- **JUST RUN IT!!!**
- Don't be afraid:
  - Running a bugged code will not harm your computer!
  - (it also should not hurt your self-esteem)
  - It doesn't cost anything
  - It will be faster (and more accurate) than you trying to "think it through"
  - In many cases, the error message or output will be extremely informative

***"The freedom to run experiments is the most precious luxury of computational biologists"***

**Nanahle Nietsnerob**

# Sample problem #1

- Download the course webpage (e.g., use the “save as” option). Write a program that reads this webpage text and scan for all the email addresses in it.
- An email address usually follows these guidelines:
  - Upper or lower case letters or digits
  - Starting with a letter
  - Followed by a the “@” symbol
  - Followed by a string of alphanumeric characters. No spaces are allowed
  - Followed by a the dot “.” symbol
  - Followed by a domain extension. Assume domain extensions are always 3 alphanumeric characters long (e.g., “com”, “edu”, “net”).

# Solution #1

```
import sys
import re

file_name = sys.argv[1]
file = open(file_name, "r")
text = file.read()

addresses = re.findall(r'[a-zA-Z]\w*@\w+\.\w{3,3}', text)
print addresses
```

```
['jht@uw.edu', 'elbo@uw.edu']
```

# Sample problem #2

1. Download and save warandpeace.txt. Write a program to read it line-by-line. Use re.findall to check whether the current line contains one or more “proper” names ending in “...ski”. If so, print these names:

```
['Bolkonski']  
['Bolkonski']  
['Bolkonski']  
['Bolkonski']  
['Volkonski']  
['Volkonski']  
['Volkonski']
```

2. Now, instead of printing these names for each line, insert them into a dictionary and just print all the “...ski” names that appear in the text at the end of your program (preferably sorted):

```
Aski  
Bitski  
Bolkonski  
Borovitski  
Bronnitski  
Czartoryski  
Golukhovski  
Gruzinski
```

# Solution #2.1

```
import sys
import re

file_name = sys.argv[1]
file = open(file_name, "r")

names_dict = {} # A dictionary for storing all names
for line in file:
    names = re.findall(r'\w+ski', line)
    if len(names) > 0:
        print names

file.close()
```

# Solution #2.2

```
import sys
import re

file_name = sys.argv[1]
file = open(file_name, "r")

names_dict = {} # A dictionary for storing all names
for line in file:
    names = re.findall(r'\w+ski', line)
    for name in names:
        names_dict[name] = 1

file.close()

name_list = names_dict.keys()
name_list.sort()

for name in name_list:
    print name
```

# Challenge problem

- “Translate” War and Peace to Pig Latin.
- The rules of translations are as follows:
  - If a word starts with a consonant: move it to the end and append “ay”
  - Else, for words that starts with a vowel, keep as is, but add “zay” at the end
  - Examples:
    - beast → eastbay
    - dough → oughday
    - happy → appyhay
    - another → anotherzay
    - if → ifzay

